

RSM8512 Assignment - Classification

Yanbing Chen

1009958752

2023/10/19

Question 3 [20 marks] (Hint: It might be easier to solve Question #2 from ISL first)

This problem relates to the QDA model, in which the observations within each class are drawn from a normal distribution with a class- specific mean vector and a class specific covariance matrix. We consider the simple case where $p = 1$; i.e. there is only one feature.

Suppose that we have K classes, and that if an observation belongs to the k th class then X comes from a one-dimensional normal distribution, $X \sim N(\mu_k, \sigma_k^2)$. Recall that the density function for the one-dimensional normal distribution is given in (4.11). Prove that in this case, the Bayes' classifier is not linear. Argue that it is in fact quadratic.

Hint: For this problem, you should follow the arguments laid out in Section 4.4.2, but without making the assumption that $\sigma_1^2 = \dots = \sigma_K^2$.

Answer: We have the Bayes Classifier is $P_r(Y = k|X = x) = \frac{\pi_k f_k(x)}{\sum_{l=1}^K \pi_l f_l(x)}$ as equation (1) and normal density function $f_k(x) = \frac{1}{\sqrt{2\pi}\sigma_k} \exp(-\frac{1}{2\sigma_k^2}(x - \mu_k)^2)$ as equation(2). By plugging (2) into (1), we have equation (3):

$$p_k(x) = \frac{\pi_k \frac{1}{\sqrt{2\pi}\sigma_k} \exp(-\frac{1}{2\sigma_k^2}(x - \mu_k)^2)}{\sum_{l=1}^K \pi_l \frac{1}{\sqrt{2\pi}\sigma_l} \exp(-\frac{1}{2\sigma_l^2}(x - \mu_l)^2)}$$

For QDA model, it assumes that each class has its own covariance matrix, so $\sigma_1 \neq \sigma_2 \neq \dots \neq \sigma_k$. After taking log of (3) and rearranging the terms and rearranging terms, we have

$$\log p_k(x) = \frac{\log(\pi_k) + \log(\frac{1}{\sqrt{2\pi}\sigma_k}) + (-\frac{1}{2\sigma_k^2}(x - \mu_k)^2)}{\log(\sum_{l=1}^K \pi_l \frac{1}{\sqrt{2\pi}\sigma_l} \exp(-\frac{1}{2\sigma_l^2}(x - \mu_l)^2))}$$

Since $\log(\sum_{l=1}^K \pi_l \frac{1}{\sqrt{2\pi}\sigma_l} \exp(-\frac{1}{2\sigma_l^2}(x - \mu_l)^2))$ is a constant, and it will not make effect on the function, therefore we have:

$$\delta_k(x) = \log(\pi_k) + \log(\frac{1}{\sqrt{2\pi}\sigma_k}) + (-\frac{1}{2\sigma_k^2}(x - \mu_k)^2)$$

As you can see from the equation above, $\delta_k(x)$ is a quadratic function of x , so the Bayes' classifier is not linear.

Question 4 [20marks] a-d from ISL. (5 BONUS marks: 4e)

When the number of features p is large, there tends to be a deterioration in the performance of KNN and other local approaches that perform prediction using only observations that are near the test observation for which a prediction must be made. This phenomenon is known as the curse of dimensionality, and it ties into the fact that non-parametric approaches often perform poorly when p is large. We will now investigate this curse.

- (a) Suppose that we have a set of observations, each with measurements on $p = 1$ feature, X . We assume that X is uniformly (evenly) distributed on $[0,1]$. Associated with each observation is a response value. Suppose that we wish to predict a test observation's response using only observations that are within 10% of the range of X closest to that test observation. For instance, in order to predict the response for a test observation with $X = 0.6$, we will use observations in the range $[0.55, 0.65]$. On average, what fraction of the available observations will we use to make the prediction?

Answer: When we do not include the situation of $X > 0.95$ and $X < 0.05$, the answer is average 10%.

- (b) Now suppose that we have a set of observations, each with measurements on $p = 2$ features, X_1 and X_2 . We assume that (X_1, X_2) are uniformly distributed on $[0,1] \times [0,1]$. We wish to predict a test observation's response using only observations that are within 10 % of the range of X_1 and within 10 % of the range of X_2 closest to that test observation. For instance, in order to predict the response for a test observation with $X_1 = 0.6$ and $X_2 = 0.35$, we will use observations in the range $[0.55, 0.65]$ for X_1 and in the range $[0.3, 0.4]$ for X_2 . On average, what fraction of the available observations will we use to make the prediction?

Answer: The fraction of the available observation is $(10\% * 10\%)/(1 * 1) = 10\%^2$

- (c) Now suppose that we have a set of observations on $p = 100$ features. Again the observations are uniformly distributed on each feature, and again each feature ranges in value from 0 to 1. We wish to predict a test observation's response using observations within the 10 % of each feature's range that is closest to that test observation. What fraction of the available observations will we use to make the prediction?

Answer: $(10\%^{100})/(1^{100}) = 10^{-100}$

- (d) Using your answers to parts (a)–(c), argue that a drawback of KNN when p is large is that there are very few training observations “near” any given test observation.

Answer: When p is increased linear, the fraction is decreased exponentially.

- (e) Now suppose that we wish to make a prediction for a test observation by creating a p-dimensional hypercube centered around the test observation that contains, on average, 10 % of the training observations. For $p = 1, 2$, and 100, what is the length of each side of the hypercube? Comment on your answer.

Answer: average fraction = $10\%^p$

Question 6 [20 marks]

Suppose we collect data for a group of students in a statistics class with variables X_1 = hours studied, X_2 = undergrad GPA, and Y = receive an A. We fit a logistic regression and produce estimated coefficient, $\hat{\beta}_0 = -6, \hat{\beta}_1 = 0.05, \hat{\beta}_2 = 1$

The probability function is

$$\hat{p}(X) = \frac{e^{\hat{\beta}_0 + \hat{\beta}_1 X_1 + \hat{\beta}_2 X_2}}{1 + e^{\hat{\beta}_0 + \hat{\beta}_1 X_1 + \hat{\beta}_2 X_2}}$$

- (a) Estimate the probability that a student who studies for 40 h and has an undergrad GPA of 3.5 gets an A in the class.

```
beta_0 = -6
beta_1 = 0.05
beta_2 = 1
x1 = 40
x2 = 3.5
exp(beta_0 + beta_1*x1+ beta_2*x2)/(1 + exp(beta_0 + beta_1*x1 + beta_2*x2))
```

```
## [1] 0.3775407
```

- (b) How many hours would the student in part (a) need to study to have a 50 % chance of getting an A in the class?

Answer: After transformation, we can get $e^{\hat{\beta}_0 + \hat{\beta}_1 X_1 + \hat{\beta}_2 X_2} = 1$, Therefore, from $\hat{\beta}_0 + \hat{\beta}_1 X_1 + \hat{\beta}_2 X_2 = 0$ we can have $-6 + 0.05x_1 + 3.5 * 1 = 0$

```
(0 - beta_0 - beta_2 * x2) / beta_1
```

```
## [1] 50
```

The student need to study 50 hours to have 50% chance of getting an A.

Question 14 [40 marks]

In this problem, you will develop a model to predict whether a given car gets high or low gas mileage based on the Auto data set.

- (a) Create a binary variable, `mpg01`, that contains a 1 if `mpg` contains a value above its median, and a 0 if `mpg` contains a value below its median. You can compute the median using the `median()` function. Note you may find it helpful to use the `data.frame()` function to create a single data set containing both `mpg01` and the other Auto variables.

```
data(Auto)
mpg01 = ifelse(Auto$mpg > median(Auto$mpg), 1, 0)
Auto = data.frame(Auto, mpg01)
#head(Auto)
```

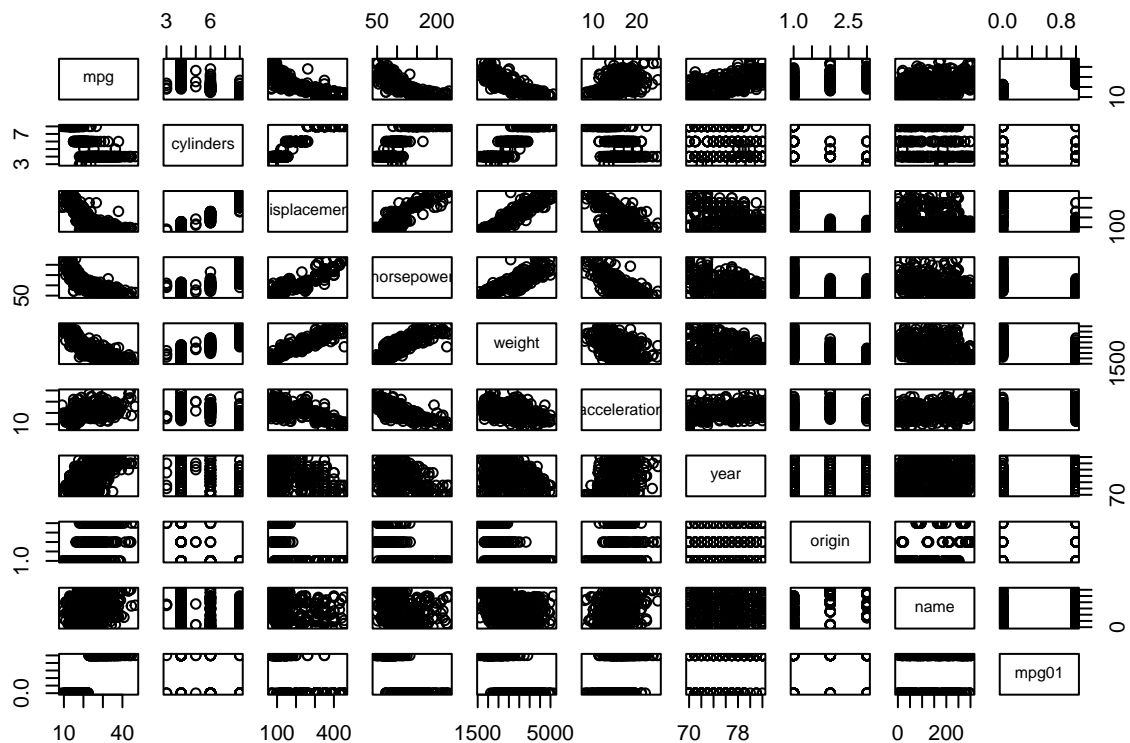
- (b) Explore the data graphically in order to investigate the association between `mpg01` and the other features. Which of the other features seem most likely to be useful in predicting `mpg01`? Scatterplots and boxplots may be useful tools to answer this question. Describe your findings.

```
cor(Auto[, -9])
```

```
##           mpg  cylinders displacement horsepower      weight
## mpg          1.0000000 -0.7776175   -0.8051269 -0.7784268 -0.8322442
## cylinders    -0.7776175  1.0000000    0.9508233  0.8429834  0.8975273
## displacement -0.8051269  0.9508233    1.0000000  0.8972570  0.9329944
## horsepower   -0.7784268  0.8429834    0.8972570  1.0000000  0.8645377
## weight       -0.8322442  0.8975273    0.9329944  0.8645377  1.0000000
## acceleration  0.4233285 -0.5046834   -0.5438005 -0.6891955 -0.4168392
## year          0.5805410 -0.3456474   -0.3698552 -0.4163615 -0.3091199
## origin        0.5652088 -0.5689316   -0.6145351 -0.4551715 -0.5850054
## mpg01         0.8369392 -0.7591939   -0.7534766 -0.6670526 -0.7577566
##
## acceleration      year      origin      mpg01
## mpg              0.4233285  0.5805410  0.5652088  0.8369392
## cylinders        -0.5046834 -0.3456474 -0.5689316 -0.7591939
## displacement     -0.5438005 -0.3698552 -0.6145351 -0.7534766
## horsepower       -0.6891955 -0.4163615 -0.4551715 -0.6670526
## weight           -0.4168392 -0.3091199 -0.5850054 -0.7577566
## acceleration     1.0000000  0.2903161  0.2127458  0.3468215
```

```
## year          0.2903161  1.0000000  0.1815277  0.4299042
## origin        0.2127458  0.1815277  1.0000000  0.5136984
## mpg01         0.3468215  0.4299042  0.5136984  1.0000000
```

```
pairs(Auto)
```



Answer: displacement, horsepower, weight and cylinders are good to fit the mpg01

(c) Split the data into a training set and a test set.

```
set.seed(1234)

split = sample.split(Auto$mpg01, SplitRatio = 0.7)
train_set = subset(Auto, split == TRUE)
test_set = subset(Auto, split == FALSE)
```

(d) Perform LDA on the training data in order to predict mpg01 using the variables that seemed most associated with mpg01 in (b). What is the test error of the model obtained?

```
lda.fit = lda(mpg01 ~ displacement + horsepower + weight + cylinders, data = train_set)

lda.pred = predict(lda.fit, test_set)
#lda.pred$class
lda_ter = 1 - mean(lda.pred$class == test_set$mpg01)
lda_ter

## [1] 0.09322034
```

(e) Perform QDA on the training data in order to predict mpg01 using the variables that seemed most associated with mpg01 in(b). What is the test error of the model obtained?

```
qda.fit = qda(mpg01 ~ displacement + horsepower + weight + cylinders, data = train_set)

qda.pred = predict(qda.fit, test_set)
qda_ter = 1 - mean(qda.pred$class == test_set$mpg01)
qda_ter

## [1] 0.1271186
```

(f) Perform logistic regression on the training data in order to predict mpg01 using the variables that seemed most associated with mpg01 in (b). What is the test error of the model obtained?

```
glm.fit = glm(mpg01 ~ displacement + horsepower + weight + cylinders, data = train_set, family = "binomial")
#summary(glm.fit)

glm.probs = predict(glm.fit, data = test_set, type = 'response')
glm.pred = rep(0, length(glm.probs))
glm.pred[glm.probs > 0.5] = 1
glm_ter = 1 - mean(glm.pred == test_set$mpg01)

## Warning in glm.pred == test_set$mpg01: longer object length is not a multiple of
## shorter object length

glm_ter # logistic regression model's test_error

## [1] 0.4963504
```

- (g) Perform naive Bayes on the training data in order to predict mpg01 using the variables that seemed most associated with mpg01 in (b). What is the test error of the model obtained?

```
library(e1071)
nb.fit = naiveBayes(mpg01 ~ displacement + horsepower + weight + cylinders, data = train_set)
nb.pred = predict(nb.fit, test_set)
nb_ter = 1 - mean(nb.pred == test_set$mpg01)
nb_ter # naive Bayes's test_error
```

```
## [1] 0.1101695
```

- (h) Perform KNN on the training data, with several values of K, in order to predict mpg01. Use only the variables that seemed most associated with mpg01 in (b). What test errors do you obtain? Which value of K seems to perform the best on this data set?

```
train.X = train_set[, c('cylinders', 'weight', 'displacement', 'horsepower')]
test.X = train_set[, c('cylinders', 'weight', 'displacement', 'horsepower')]
train_mpg01 = train_set[, 'mpg01']
test_mpg01 = test_set[, 'mpg01']

set.seed(1234)
```

```
# when k = 1
knn.pred = knn(train.X, test.X, train_mpg01, test_mpg01, k = 1)
knn_ter1 = 1 - mean(knn.pred == test_mpg01 )
```

```
## Warning in `==.default`(knn.pred, test_mpg01): longer object length is not a
## multiple of shorter object length
```

```
## Warning in is.na(e1) | is.na(e2): longer object length is not a multiple of
## shorter object length
```

```
knn_ter1
```

```
## [1] 0.4817518
```

```
# when k = 10
```

```
knn.pred = knn(train.X, test.X, train_mpg01, test_mpg01, k = 10)
```

```
knn_ter10 = 1 - mean(knn.pred == test_mpg01 )
```

```
## Warning in `==.default`(knn.pred, test_mpg01): longer object length is not a  
## multiple of shorter object length
```

```
## Warning in is.na(e1) | is.na(e2): longer object length is not a multiple of  
## shorter object length
```

```
knn_ter10
```

```
## [1] 0.5072993
```

```
# when k = 100
```

```
knn.pred = knn(train.X, test.X, train_mpg01, test_mpg01, k = 100)
```

```
knn_ter100 = 1 - mean(knn.pred == test_mpg01 )
```

```
## Warning in `==.default`(knn.pred, test_mpg01): longer object length is not a  
## multiple of shorter object length
```

```
## Warning in is.na(e1) | is.na(e2): longer object length is not a multiple of  
## shorter object length
```

```
knn_ter100
```

```
## [1] 0.4963504
```